

Introduction to Machine Learning Applications

Spring 2023

Deep learning

Minor Gordon

gordom6@rpi.edu




Rensselaer

Deep Learning

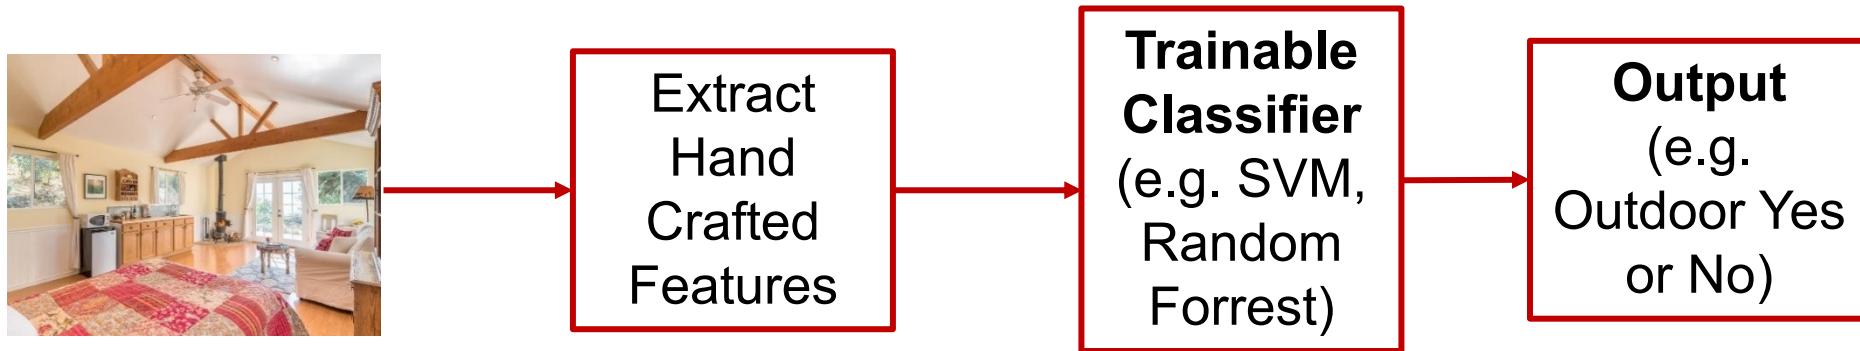
- Neural networks are universal function approximators.
- Deep learning combines large neural network architectures along with innovative training algorithms (Convolutions, RNN, GANS, Autoencoders, etc.) along with (typically) large datasets to build predictive models
- Because of size of data and complexities of operations dedicated hardware (GPU) is often required

Performance vs Sample Size



Traditional Supervised Learning

- Traditional pattern recognition models work with hand crafted features and relatively simple trainable classifiers.





Limitations


- Very tedious and costly to develop hand-crafted features
- The hand-crafted features are usually highly dependent on one application.

Deep Learning

- Deep learning has a **built-in automatic multi-stage feature learning process** that learns rich hierarchical representations (i.e., features).



Deep Learning



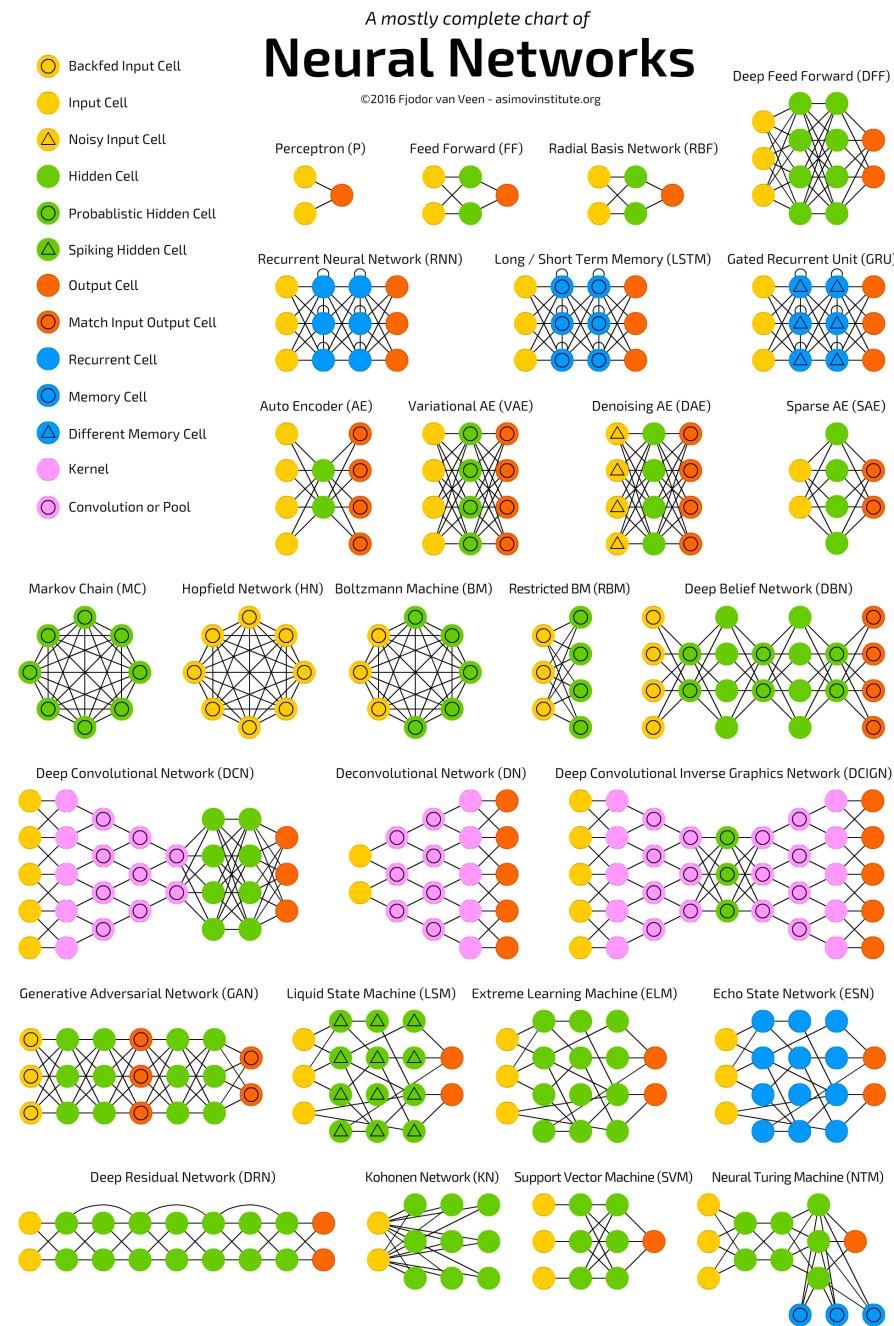
- Image

Pixel → Edge → Texture → Motif → Part → Object

- Text

Character → Word → Word-group → Clause → Sentence → Story

- Each module in Deep Learning transforms its input representation into a higher-level one, in a way similar to human cortex.



Many
different
types of
neural
networks,
with types
based on
type of
problem

Source: <http://www.asimovinstitute.org/neural-network-zoo/>

Supervised Convolutional Neural Network

Why Convolutional Neural Networks?

A CNN is a deep, feed-forward artificial neural network that has successfully been applied to analyzing visual imagery.

ImageNet Challenge


IMAGENET

- 1,000 object classes (categories).
- Images:
 - 1.2 M train
 - 100k test.



Convolutional Networks

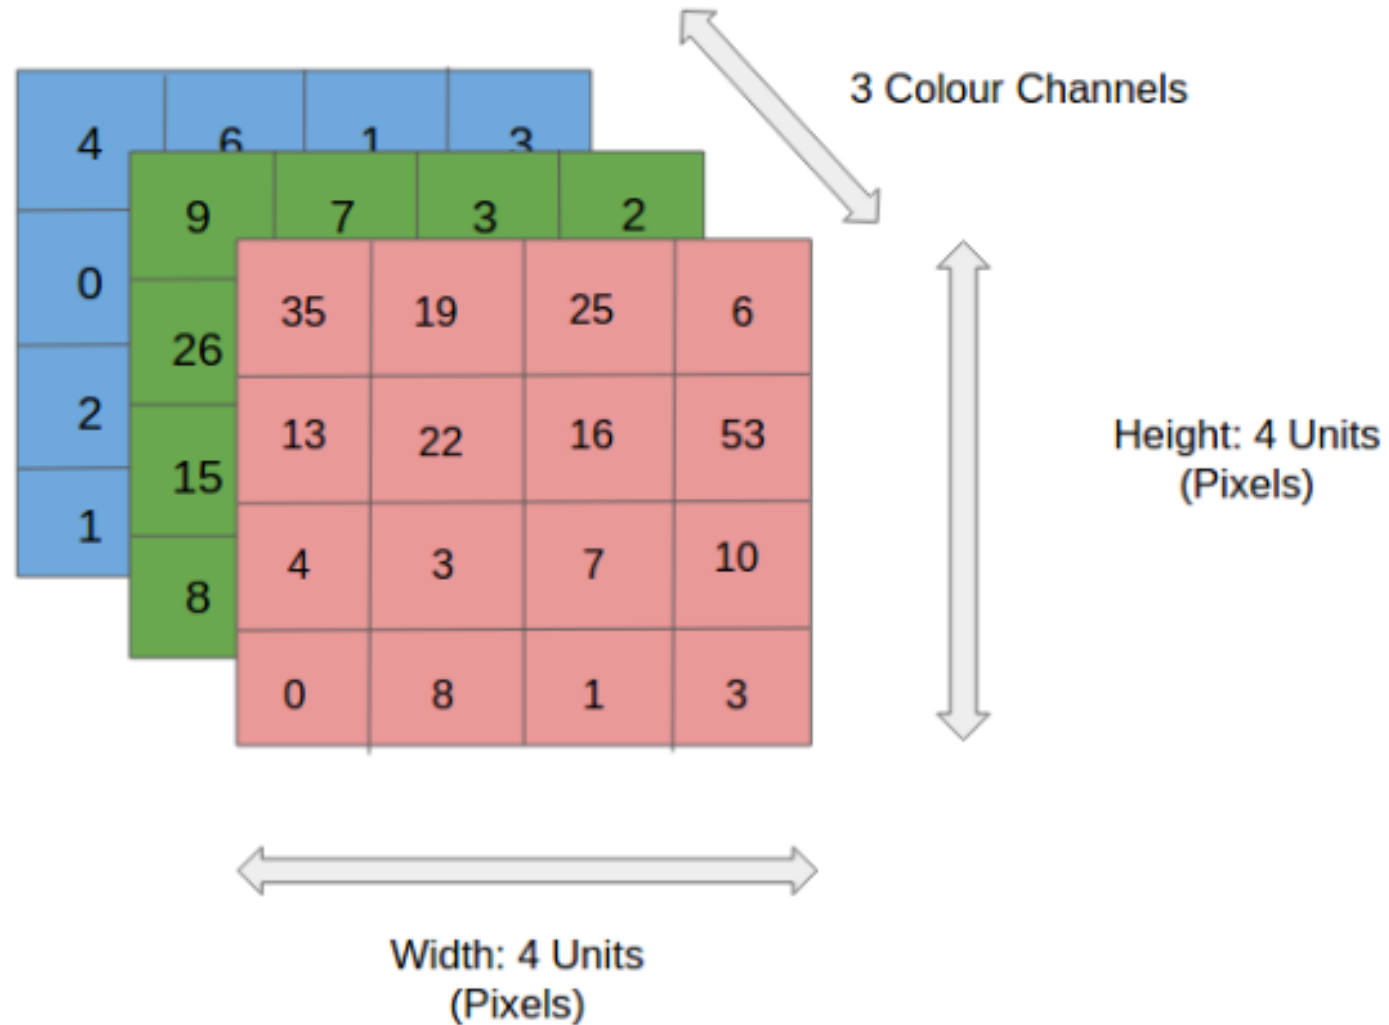
- Image processing
- Sentence Classification




Convolutional Neural Networks

- **Output:** Binary, Multinomial, Continuous, Count
- **Input:** fixed size, can use padding to make all images same size.
- **Architecture:** Choice is ad hoc
 - requires experimentation.
- **Optimization:** Backward propagation
 - hyper parameters for very deep model can be estimated properly only if you have billions of images.
 - Use an architecture and trained hyper parameters from other papers (Imagenet or Microsoft/Google APIs etc)

Input Image



Input an Image to a Neural Network



Number of Parameters

$480000 \times 480000 + 480000 + 1 = \text{approximately } 230 \text{ Billion !!!}$

$480000 \times 1000 + 1000 + 1 = \text{approximately } 480 \text{ million !!!}$

Convolutional Layers

- Filter 
$$\begin{pmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{pmatrix}$$



Input Image

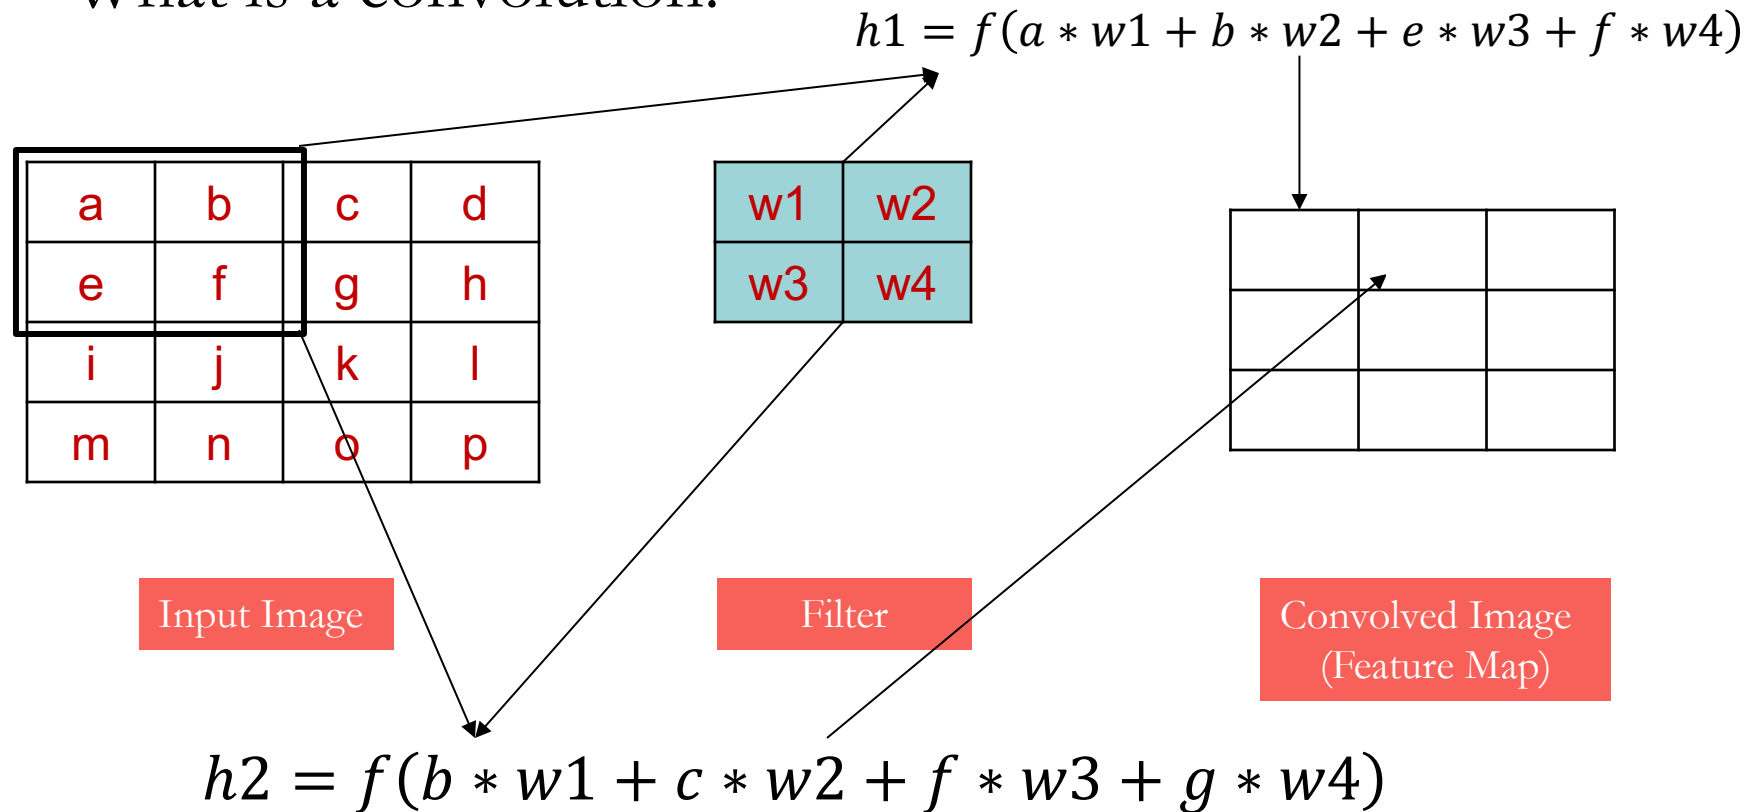


Convolved Image

- Inspired by the neurophysiological experiments conducted by Hubel and Wiesel 1962.

Convolutional Layers

- What is a convolution?



Number of Parameters for one feature map = 4

Number of Parameters for 100 feature map = 4*100

Convolution Layer – The Kernel

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

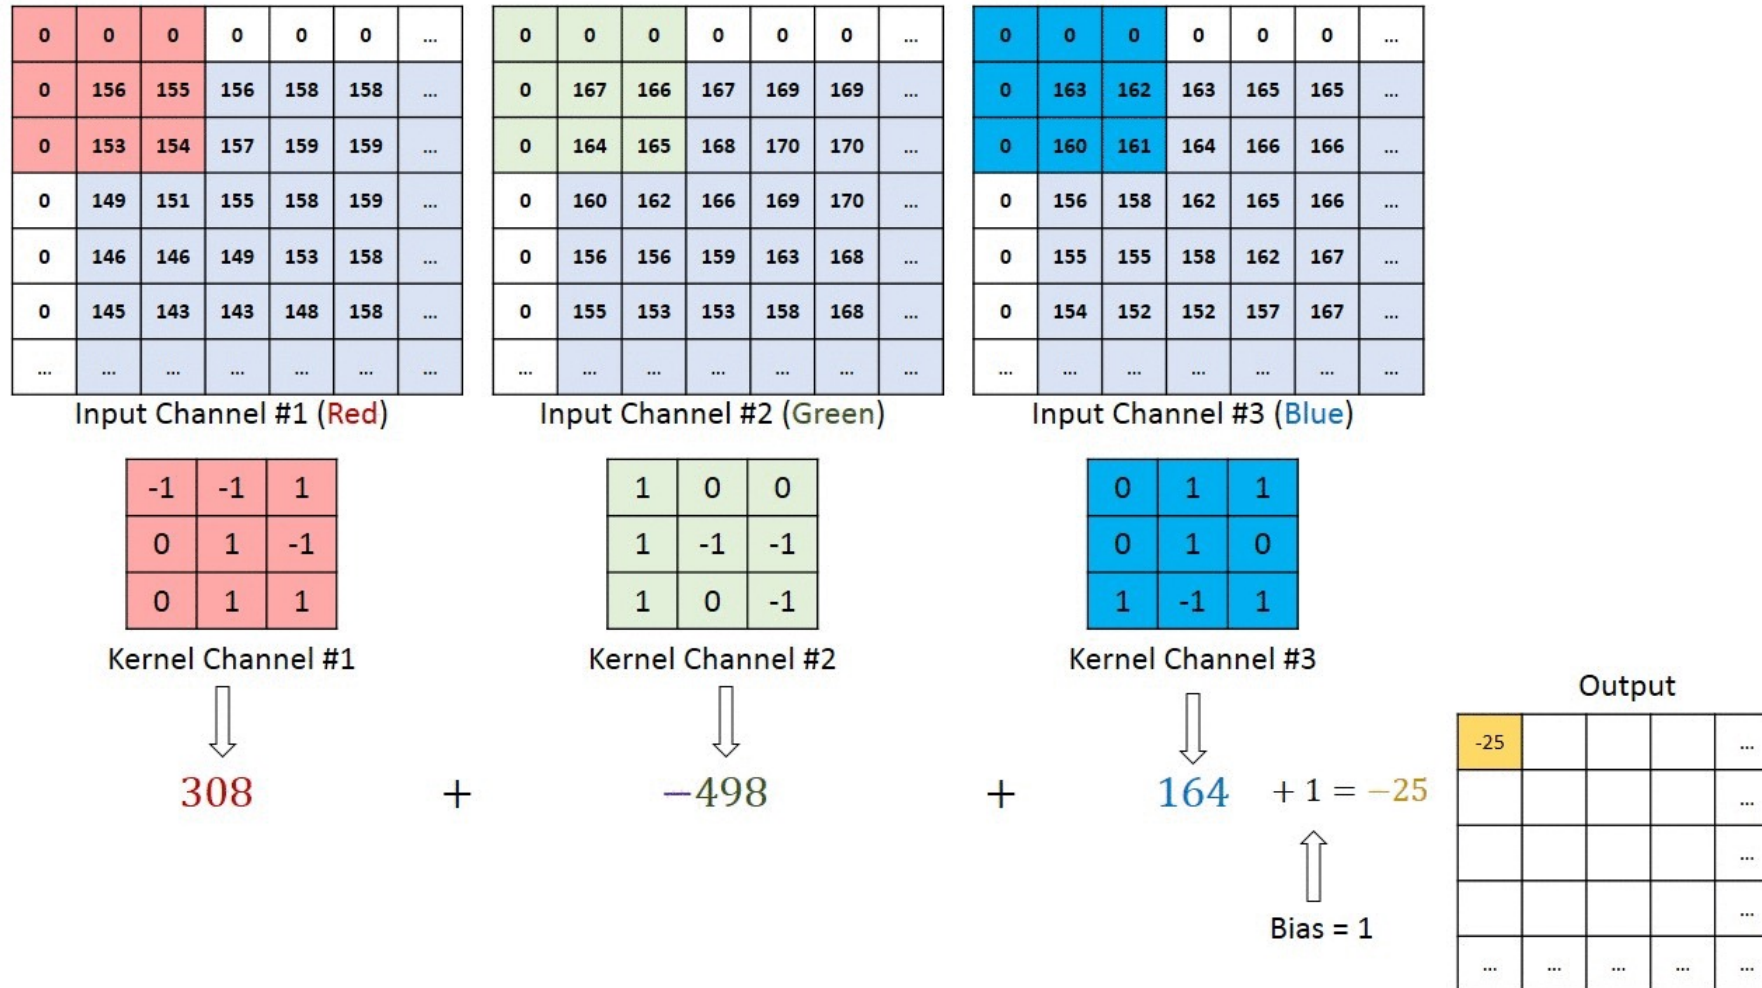
Kernel/Filter, K =

1 0 1

0 1 0

1 0 1

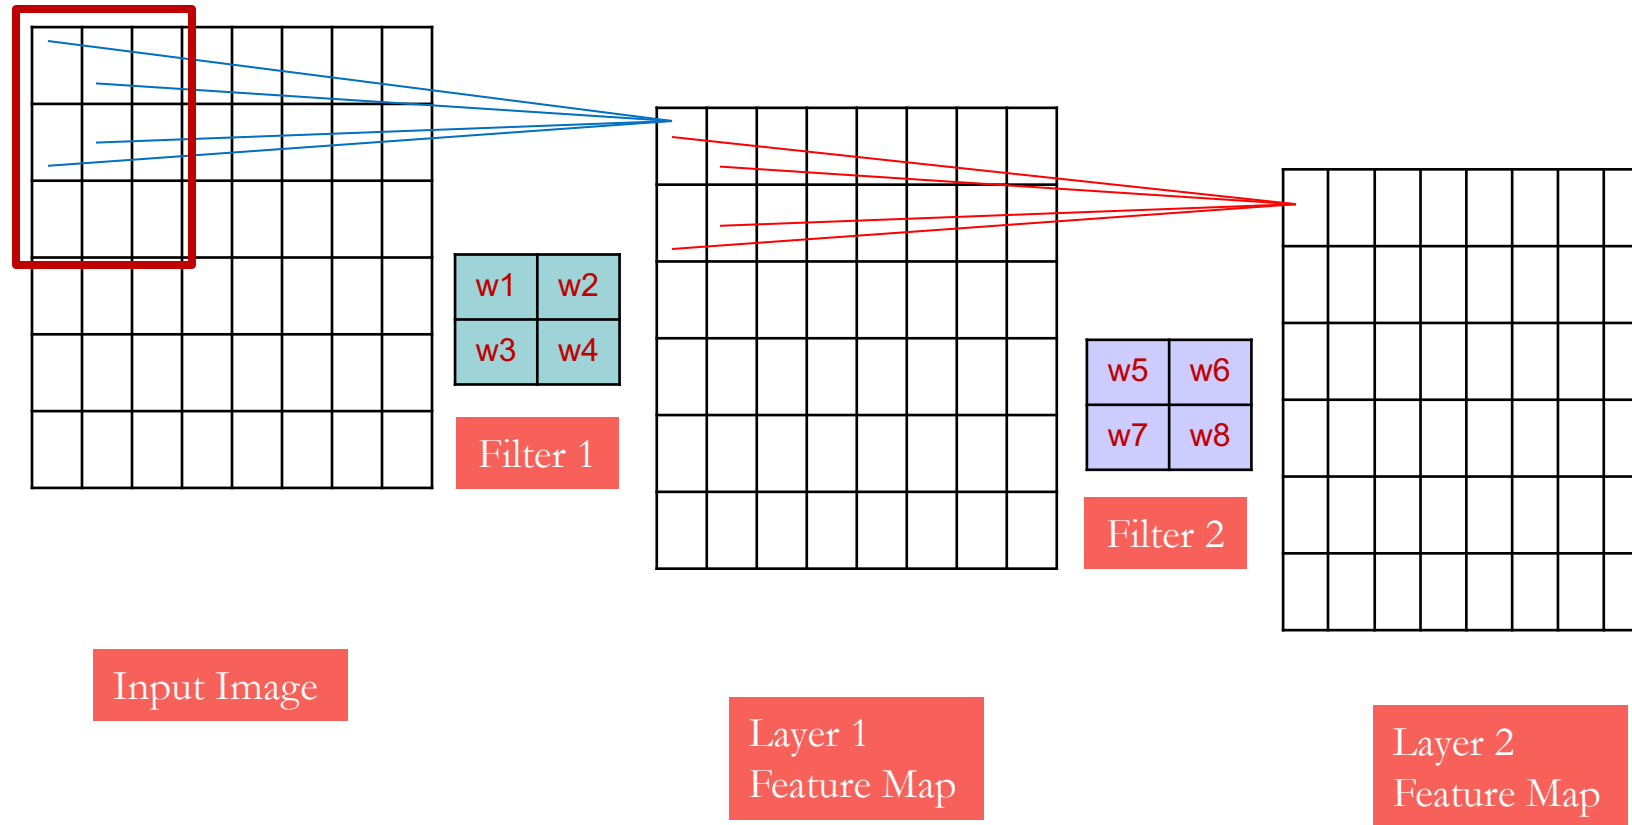
Convolution operation on a MxNx3 image



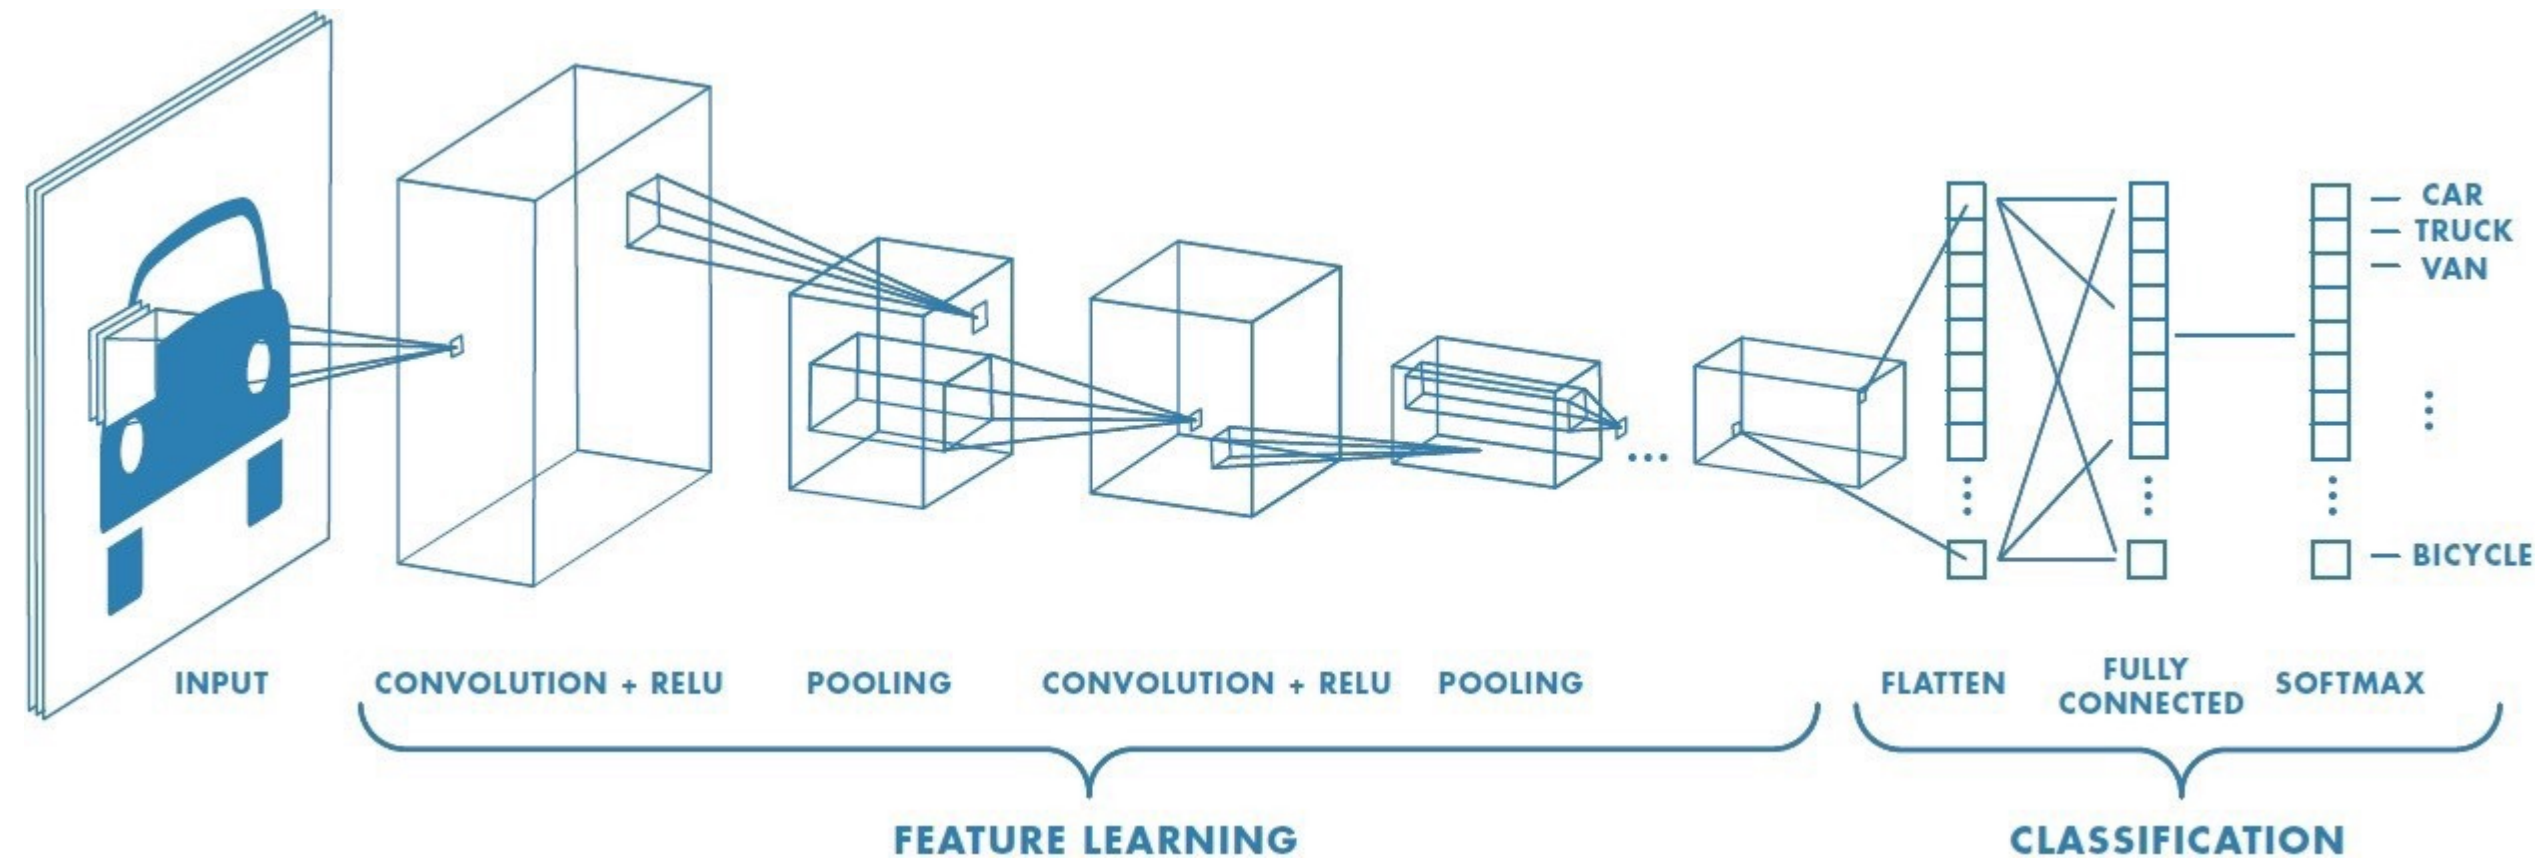
Convolution operation

- The objective of the Convolution Operation is to **extract the high-level features** such as edges, from the input image.
- ConvNets need not be limited to only one Convolutional Layer.
- Conventionally, the first ConvLayer is responsible for capturing low-level features such as edges, color, gradient orientation, etc.
- With added layers, the architecture adapts to high-level features as well.

Lower Level to More Complex Features



- In Convolutional neural networks, hidden units are only connected to local receptive field.



Pooling

- **Max pooling:** reports the maximum output within a rectangular neighborhood.
- **Average pooling:** reports the average output of a rectangular neighborhood.

1	3	5	3
4	2	3	1
3	1	1	3
0	1	0	4

Input Matrix

MaxPool with 2X2 filter with
stride of 2

4	5
3	4

Output Matrix

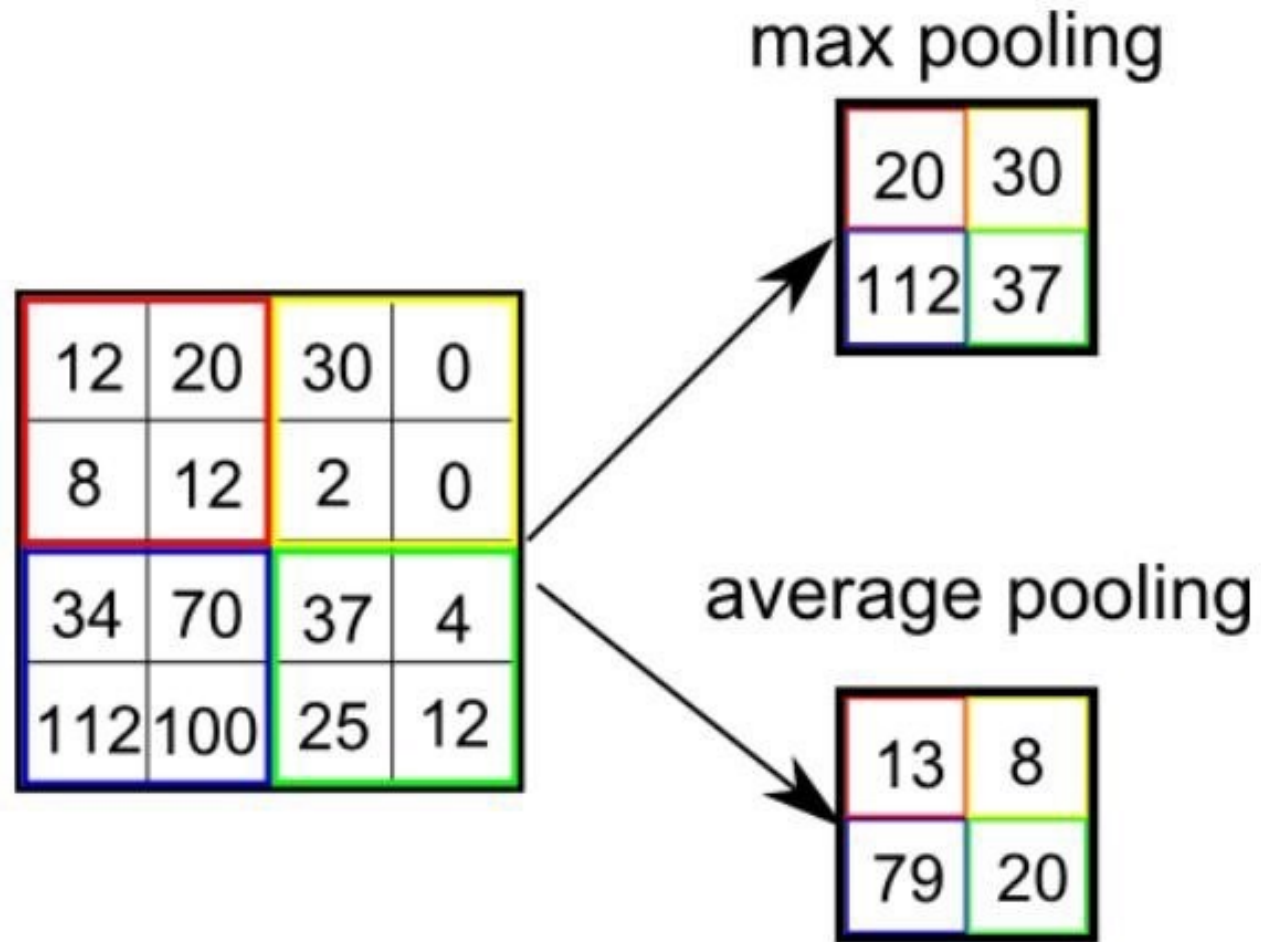
Pooling layer

3.0	3.0	3.0
3.0	3.0	3.0
3.0	2.0	3.0

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

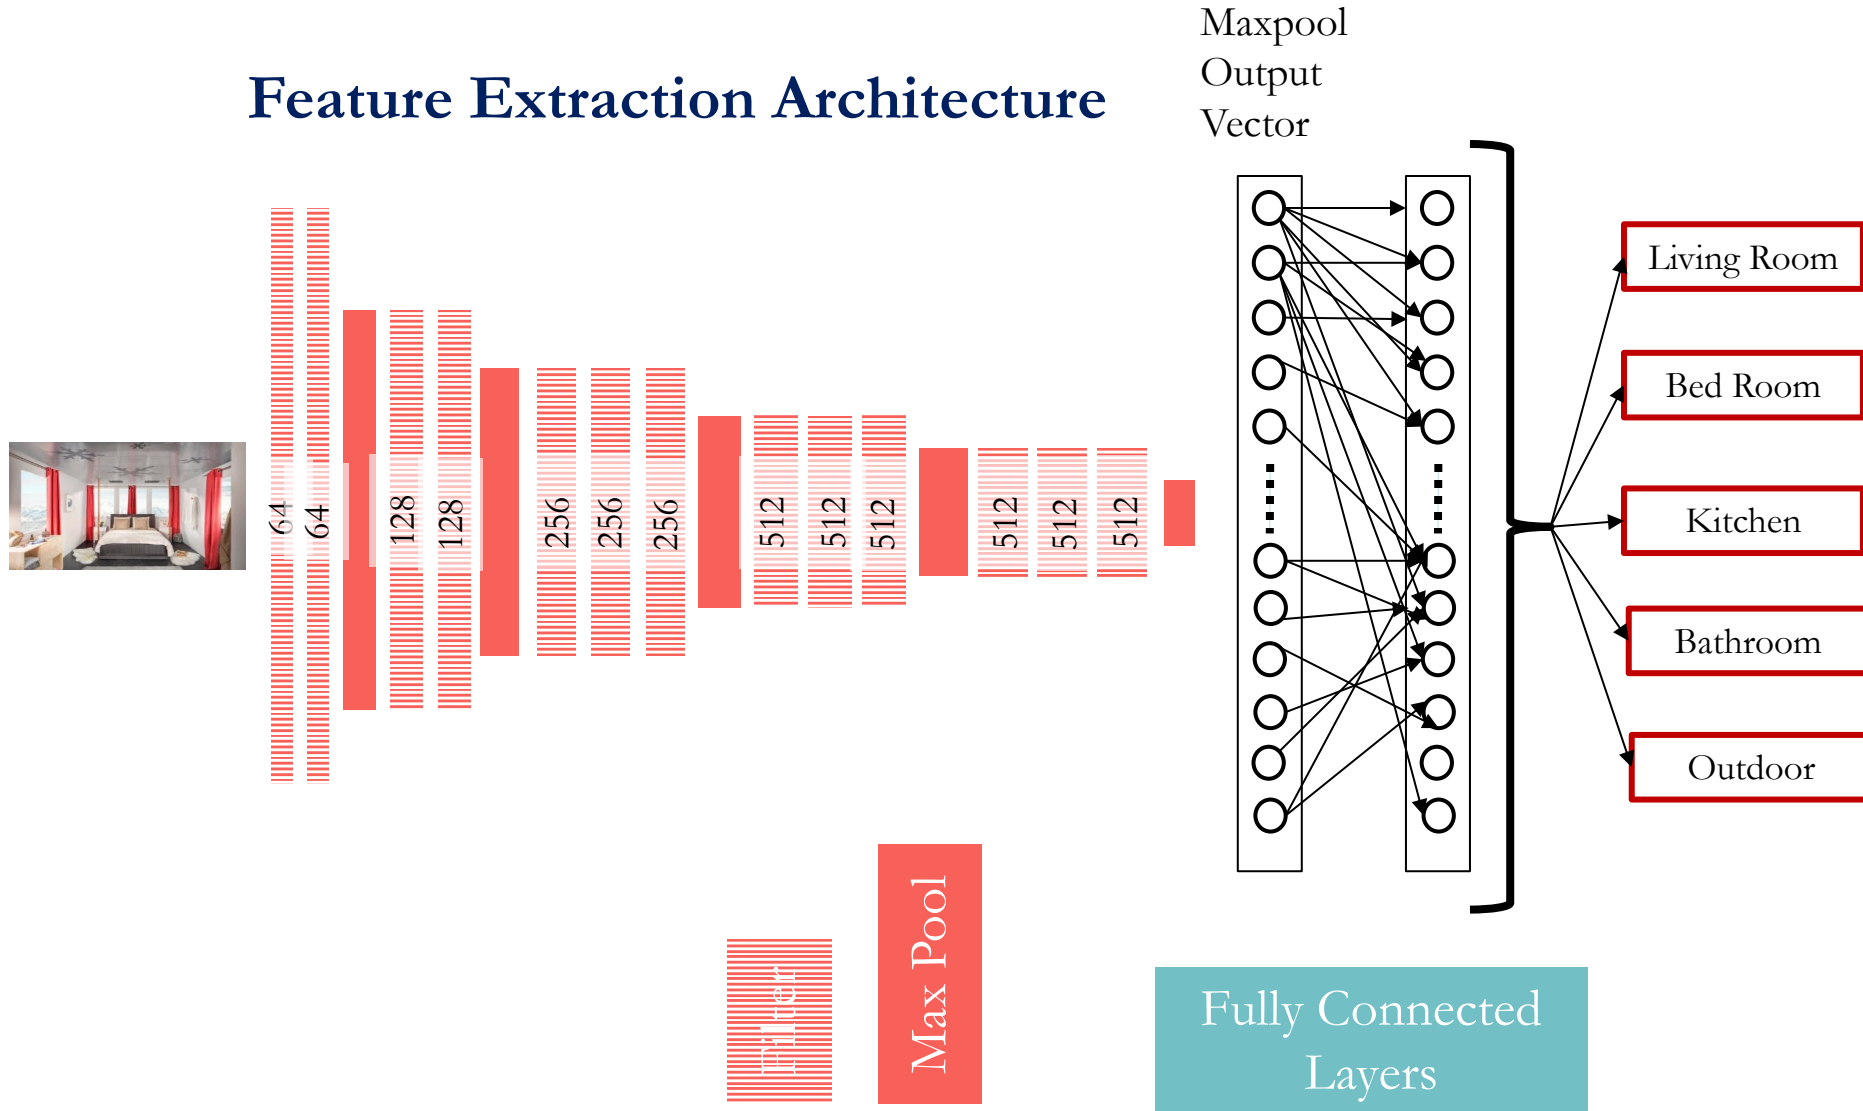
- **Max Pooling** returns the **maximum value** from the portion of the image covered by the Kernel.
- **Average Pooling** returns the **average of all the values** from the portion of the image covered by the Kernel.

Pooling example



Convolutional Neural Network

Feature Extraction Architecture



Batch normalization

- Main goal is to increase the stability of a neural network.
- This technique normalizes the output of a previous activation layer by subtracting the batch mean and dividing by the batch standard deviation.
- We can use higher learning rates because batch normalization makes sure that there's no activation that's gone really high or really low.
- Reduces overfitting
 - By using this approach, we can avoid dropout, and thus avoid losing information.

Pretrained Models and Transfer Learning

Pretrained Models – Key to a Quick Start

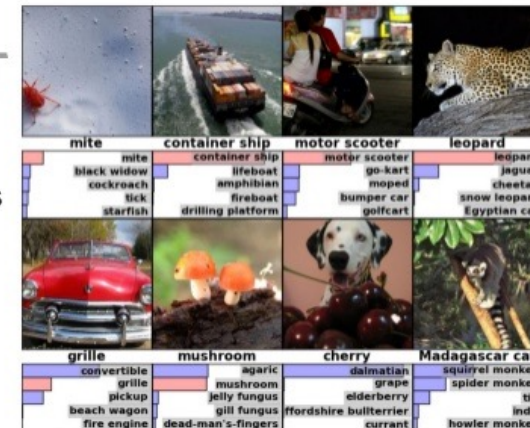
- Possible to take already trained models and retrain the last layer to do something specific, even if the model has been trained on something general.

The Inception Model:
Deep Convolutional
Neural Network
Capable of identifying
1,000 object classes
trained from 100K
images.

ImageNet Challenge

IMAGENET

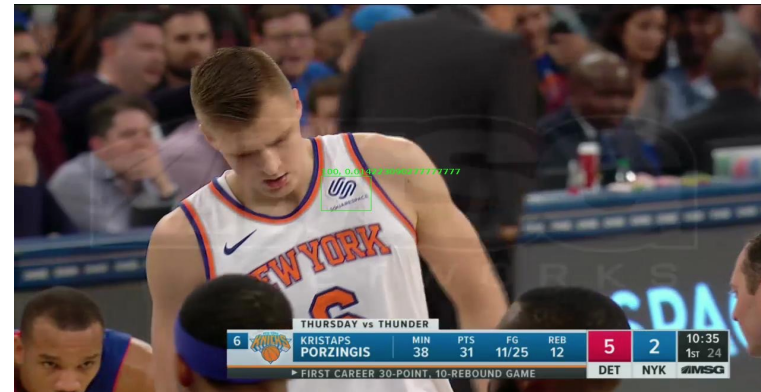
- 1,000 object classes (categories).
- Images:
 - 1.2 M train
 - 100k test.



<https://github.com/tensorflow/models>

Pretrained Models – Reusing Last Layer

- Let's say you need to identify Star Wars Characters to detect for Copyright violations.
- Or detect logos in NBA games



https://github.com/llSourcell/tensorflow_image_classifier

Transfer Learning

- Transfer learning is a machine learning method where a model developed for a task is reused as the starting point for a model on a second task.

Pre-trained Model Approach

1. Select Source Model
2. Reuse Model.
3. Tune Model (Maybe retrain last layers)